

# Deep Learning - PascalVOC Project

---

## Krishna Penukonda (1001781)

This document only contains snippets of the real code that may not work in an interpreter.

For the full code, please refer to the IPython notebook and accompanying python modules.

## Summary

---

Resnet50-based model with sigmoid outputs.

No learning rate scheduler.

Converges in ~10 epochs on PascalVOC training set.

## Hyperparameters

- Learning Rate: `1e-4`
- Optimizer: `Adam`
- Batch size: `8`
- Loss function: *Class-wise Binary Cross-Entropy*

## Metrics

- Accuracy: `0.9628`
- Mean AP: `0.7286`
- Mean Tail Accuracy: `0.9116`

## Dependencies

---

Install this project's dependencies via `pip install -r requirements.txt`

I use [PyTorch Lightning](#) to facilitate experiment logging. This library does not abstract any significant functionality of Pytorch other than the training loop (which is mostly copy-pasted code anyway).

```
from pathlib import Path
import shutil
from tqdm.notebook import tqdm

import numpy as np
from sklearn.metrics import average_precision_score
import matplotlib.pyplot as plt
from PIL import Image

import torch
import torchvision
import torch.nn.functional as F
from torch.utils.data import DataLoader
from torchvision import transforms
import pytorch_lightning as pl # Helps with logging and other stuff
```

```
from voc_data import VOCDataset, PascalVOC # Dataloader is defined here
```

# Model Training

## Multilabel Classifier

The classifier is based on Resnet50 (pre-trained on ImageNet).

The final fully-connected layer is replaced with a new randomly initialized layer that has 20 outputs rather than the default 1000, and the sigmoid function is applied to its outputs.

## Loss Function

The loss function I use is based on Pytorch's [binary\\_cross\\_entropy](#). This function is applied class-wise and averaged to obtain the loss value (see `VOCClassifier._multi_label_loss` for implementation).

## Data Pre-processing

The label for each sample is represented as a many-hot vector, with negative classes being `0` and positive ones being `1`.

During both training and inference, images are resized (shortest edge = 400px), FiveCropped and stacked, then normalized.

During training, I apply additional augmentation (random horizontal flipping and rotation).

## Model Definition

```
class VOCClassifier(pl.LightningModule):
    def __init__(self, learning_rate=1e-4):
        super().__init__()
        # Hyperparameters
        self.learning_rate = learning_rate

        # Model definition
        self.stem = torchvision.models.resnet50(pretrained=True, progress=True)
        self.stem.fc = torch.nn.Linear(2048, 20)
        self.sigmoid = torch.nn.Sigmoid()

    # FORWARD PASS
    def forward(self, x):
        bs, ncrops, c, h, w = x.size()
        x = self.stem(x.view(-1, c, h, w))
        x = x.view(bs, ncrops, -1).max(1)[0]
        return self.sigmoid(x)

    # LOSS FUNCTION
    @staticmethod
    def _multi_label_loss(pred, labels):
        loss = torch.stack([F.binary_cross_entropy(
            pred[:, i],
            labels[:, i]
        ) for i in range(labels.shape[1])]).mean()
```

```

        return loss

    # ONE TRAINING ITERATION (backward pass handled by Pytorch-lightning)
    def training_step(self, batch, batch_idx):
        image, labels, _ = batch
        pred = self.forward(image)
        loss = self._multi_label_loss(pred, labels)
        output = {"loss": loss}
        if self.log_every_n_steps and batch_idx % self.log_every_n_steps == 0:
            tensorboard_logs = {"train_loss": loss}
            output["log"] = tensorboard_logs
        return output

    # OPTIMIZER
    def configure_optimizers(self):
        optimizer = torch.optim.Adam(self.parameters(), lr=self.learning_rate)
        return optimizer

model = VOCClassifier(learning_rate=1e-4)

```

`pytorch_lightning` comes in here.

It handles the training loop and the moving of tensors to GPU(s), as well as logging of metrics for visualization in Tensorboard.

```

trainer = pl.Trainer(gpus=[0])
trainer.fit(model, train_dataloader=train_dataloader,
            val_data loaders=val_dataloader)

```

## Model Evaluation

**Accuracy:** 0.9628

**Class-wise Average Precision:**

```

[0.90154388 0.75927758 0.81515773 0.72510909 0.5123822 0.85176327 0.80504543
0.89883597 0.57333306 0.61246603 0.53352054 0.83452248 0.70566054 0.78140663
0.94029254 0.56839073 0.75188008 0.45492436 0.80734184 0.74007774]

```

**Mean Average Precision:** 0.7286

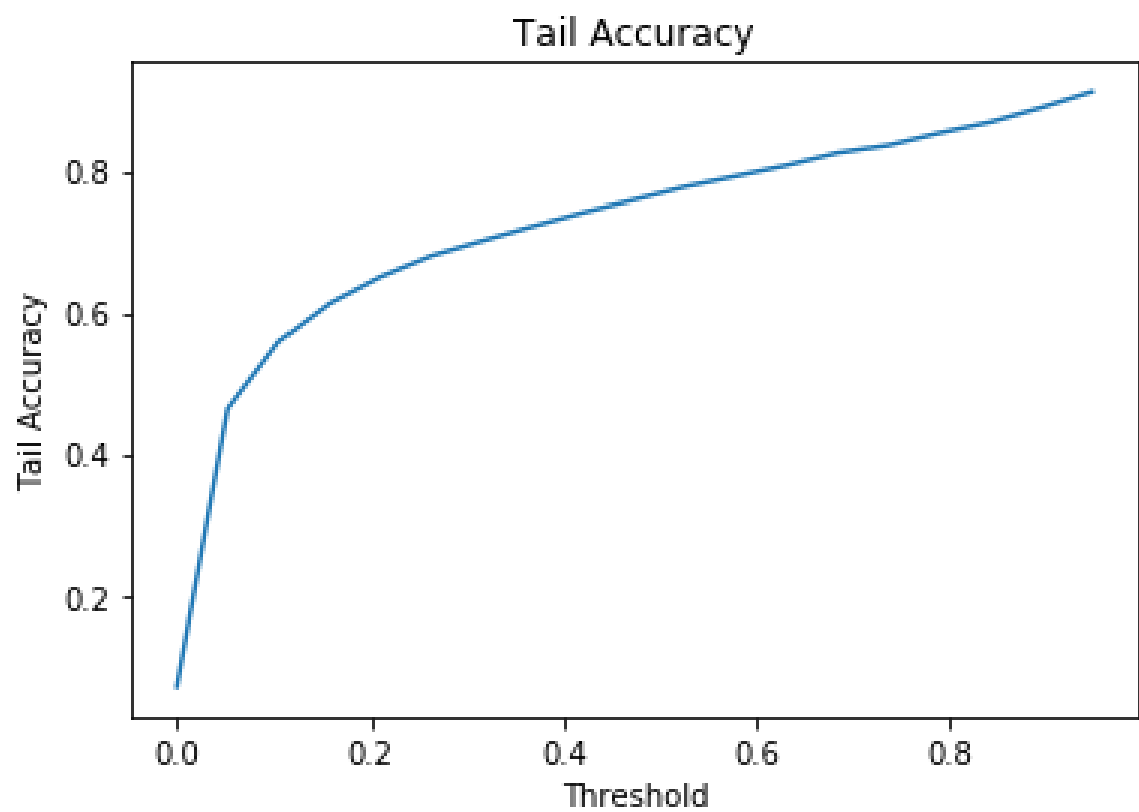
**Class-wise Tail Accuracy:**

```

[0.95522384 0.93589732 0.98275856 0.83647793 0.9342104 0.9999998 0.97435893
0.9478908 0.75555551 0.73493967 0.72093006 0.96193768 0.92941166 0.97727251
0.98116297 0.8765431 0.93589732 0.81481451 0.990909 0.98684198]

```

**Mean Tail Accuracy:** 0.9116



Class-wise Top-5 and Bottom-5 samples

cow: Top 5



cow: Bottom 5



bicycle: Top 5




bicycle: Bottom 5



car: Top 5




car: Bottom 5






person: Top 5



person: Bottom 5



tvmonitor: Top 5



tvmonitor: Bottom 5

